

Extending Equilibrium Matching: Second-Order Dynamics, Latent Spaces, and Conditional Generation

Vladislav M. Minashkin

Scientific Supervisor: Andrey V. Grabovoy, PhD

Department of Intelligent Systems
School of Applied Mathematics and Computer Science
Moscow Institute of Physics and Technology

December 2025

Equilibrium Matching

Problem Statement

Current SOTA generative models (Diffusion, Flow Matching) rely on **non-equilibrium dynamics** (time-dependent vector fields). This enforces fixed integration horizons and prevents flexible optimization-based sampling.

Research Goal

Develop extensions to Equilibrium Matching that enhance inference speed, training efficiency, and stability while rigorously preserving its inherent time-invariant vector field properties.

Tasks

- ▶ Extend EqM to conditional generation and latent spaces.
- ▶ Regularize the training using Second-Order (Hessian) dynamics.
- ▶ Analyze the sampling trajectory using UMAP.
- ▶ Analyze the stability and geometric properties of EqM.

Baseline: Equilibrium Matching

We define a corruption process via interpolation $\gamma \sim U[0, 1]$:

$$\mathbf{x}_\gamma = \gamma \mathbf{x} + (1 - \gamma) \boldsymbol{\epsilon}$$

Unlike Flow Matching, EqM minimizes the difference between the model output $f(\mathbf{x}_\gamma)$ and a target gradient field pointing from noise to data:

$$\mathcal{L}_{\text{EqM}} = \mathbb{E}_{\mathbf{x}, \boldsymbol{\epsilon}, \gamma} \left[\left\| f(\mathbf{x}_\gamma) - \underbrace{(\boldsymbol{\epsilon} - \mathbf{x})c(\gamma)}_{\text{Target Gradient}} \right\|^2 \right]$$

- ▶ **Implicit Energy:** $f(\mathbf{x}) \approx \nabla_{\mathbf{x}} E(\mathbf{x})$.
- ▶ **Optimization Sampling:** Sampling is performed via Gradient Descent (GD) rather than ODE integration.

Proposed Extension: Second-Order Regularization

Standard EqM penalizes the first derivative but ignores curvature (Hessian). We introduce a regularization term to align the curvature of the energy landscape.

Method: Using Hutchinson's Estimator for Jacobian-Vector Products (JVP). Let $\mathbf{v} \sim \mathcal{N}(0, I)$.

$$\mathcal{L}_{2\text{nd}} = \left\| \underbrace{(\nabla_{\mathbf{x}_\gamma} f(\mathbf{x}_\gamma)) \cdot \mathbf{v}}_{\text{Predicted Curvature}} - \underbrace{(-c(\gamma)\mathbf{v})}_{\text{Target Curvature}} \right\|^2$$

Hypothesis

Explicitly controlling the Hessian eigenvalues ensures the energy landscape is convex near the data manifold, improving convergence of Newton/NAG methods.

Proposed Extension: Conditional Generation

To enable class-conditional generation (e.g., MNIST/CIFAR classes), we introduce Adaptive Group Normalization (**AdaGN**).

Given class embedding $\mathbf{e}_y$:

$$\text{AdaGN}(\mathbf{h}, y) = \text{GroupNorm}(\mathbf{h}) \cdot (1 + \mathbf{s}(\mathbf{e}_y)) + \mathbf{b}(\mathbf{e}_y)$$

This modulates the energy landscape:

$$E(\mathbf{x}, y) \neq E(\mathbf{x})$$

effectively creating distinct "points of attraction" for different classes within the same vector field.

Proposed Extension: Importance Sampling & Smoothing

1. **Smooth Gradient Magnitude** $c(\gamma)$: Original EqM uses piecewise linear functions. We propose a smooth approximation to ensure differentiability of the target field:

$$c_{\text{smooth}}(\gamma) = \frac{r(\gamma)}{\sqrt{1 + r(\gamma)^2}}, \quad r(\gamma) = \frac{1 - \gamma}{1 - a}$$

2. **Gamma Scheduling**: Uniform sampling $\gamma \sim U[0, 1]$ is probably is not the most efficient. We investigate Beta Distribution sampling $\gamma \sim \text{Beta}(\alpha, \beta)$ to prioritize the "confusion zone" ($\gamma \approx 0.5$) where the gradient direction is most critical.

Experiments: Conditional Generation

We evaluated the ability of the model to generate class-specific samples using the proposed AdaGN modulation.




Figure: Qualitative results: Generated samples sorted by class.

- **Observations:** The model successfully separates class manifolds.
- **Diversity:** High intra-class variance is maintained without the need for additional classifier guidance.

Experiments: Second-Order Regularization

Initial tests with Hessian-based regularization $\mathcal{L}_{2\text{nd}}$ show promise in stabilizing the energy landscape, though the trade-off with reconstruction quality is sensitive.

- ▶ **Coefficient Fine-tuning:** The balance between the EqM loss and the second-order term is critical. Excessive regularization leads to "blurred" modes, while too little fails to improve Newton-method convergence.
- ▶ **Ongoing Work:** Automated scheduling of λ_{reg} during training.



Figure: Current samples generated with second-order dynamics.

Experiments: Smooth $c(\gamma)$ and Adaptive Sampling

Replacing the piecewise linear $c(\gamma)$ with a smooth function ensures a continuous gradient field, facilitating faster convergence in adaptive settings.



Baseline (Truncated)



Smooth $c(\gamma)$

Setting	Avg. Sampling Steps
Linear $c(\gamma)$	42.0
Smooth $c(\gamma)$	34.0

Conclusions

1. **Conditional Generation:** AdaGN successfully modulates the equilibrium points, allowing controllable generation without classifier guidance.
2. **Smooth Dynamics:** The proposed smooth $c(\gamma)$ function is superior to piecewise linear baselines, resulting in faster convergence and more stable sampling trajectories in adaptive settings.
3. **Second-Order Regularization:** Current experiments highlight significant challenges; the model is highly sensitive to regularization coefficients, requiring delicate fine-tuning to avoid over-constraining the vector field.

Next Steps

- ▶ Implement ICNN (Input Convex Neural Networks) to guarantee global energy convergence.
- ▶ Develop automated methods for scheduling second-order regularization coefficients during training.
- ▶ Scaling Equilibrium Matching to Latent Diffusion (LEqM) for high-resolution synthesis.